DEVENDRA REDDY YERVA

Agentic AI / GenAI Technical Lead | AI Agents & Multi-Agent Systems · RAG · LLMOps & Evals

Hyderabad, India · +91 70133 13692 · dev225.reddys@gmail.com · linkedin.com/in/dev225x · devendra.me

7 yrs	3 days → 2 hrs	500+	15+	92%
PRODUCTION AI/ML	EXEC REVIEW CYCLE	ENTERPRISE USERS	ENGINEERS LED	RAG RETRIEVAL ACC.

PROFILE

Engineer and technical lead with **7 years building production AI systems** for Fortune 500 clients across telecom, automotive, and medical imaging. Currently the technical owner of an **agentic AI analytics platform** in Verizon's CFO organization: multi-agent orchestration of Large Language Models (LLMs), Text-to-SQL over 50M+ row warehouses, and production Retrieval-Augmented Generation (RAG) serving 500+ enterprise users. Previously led **15+ CV/ML engineers** taking an enterprise annotation platform from concept to eight international clients. The same hands on architecture, delivery, the executive review, and the on-call pager.

PROFESSIONAL EXPERIENCE

Verizon

Oct 2024 – Present
Hyderabad, India

Consultant, AI / GenAI | Business Transformation (CFO Organization)

> Technical owner of the CFO org's agentic GenAI platform · **500+ enterprise users** · direct C-suite stakeholders · Claude Opus/Sonnet, GCP, BigQuery, Teradata

- **Cut executive business-case reviews from 3+ days to under 2 hours** by architecting BT OneView, a production agentic AI platform on GCP: sequential agent orchestration, tool-use routing, and structured-output parsing over the Claude API.
- **Opened 50M+ row Teradata and BigQuery datasets to non-technical stakeholders** through schema-aware Text-to-SQL (NL-to-SQL) pipelines, removing engineering from the query path entirely.
- **Returned 20+ analyst hours per week** with 8 autonomous overnight research agents that query the warehouse, synthesize findings, and deliver branded PDF/PPT briefings by morning, with a supervisor agent gating every briefing before it ships.
- **Built War Room, an adversarial multi-agent decision engine** where 18 specialist AI personas debate a proposal from opposing sides before a synthesis agent issues a stress-tested recommendation.
- **Reached 92% retrieval accuracy with full source attribution** on a production RAG stack (Pinecone, ChromaDB) using hybrid vector + keyword search, chunk-size optimization, re-ranking, and hallucination guardrails.
- **Reduced incident MTTR by 40%** through multi-agent auto-triage, with LLM agents performing classification, routing, and stakeholder notification without human dispatch.
- **Shipped 8 analytics dashboards to 500+ enterprise users**, holding sub-second response on warehouse-scale queries through a three-tier DuckDB cache.
- **Set the AI automation roadmap alongside CFO-org leadership** in weekly strategy reviews, presenting ROI, defending measurement methodology, and shaping investment priorities across product, data engineering, and finance.

L&T Technology Services

Feb 2022 – Sep 2024
Hyderabad, India

Senior Engineer, Transportation Solutions (Technical Lead)

> Led **15+ CV/ML engineers and data scientists** · **8 enterprise clients** across Europe & the Americas · 200+ daily users · 500K+ annotations/month

- **Took AnnotAI from concept to production**. The computer-vision annotation platform now serves 200+ daily users across 8 enterprise clients in automotive, medical imaging, and robotics at 500K+ annotations per month.
- **Led 15+ engineers with full people-management responsibility**, not only technical direction: objective-setting, performance reviews, hiring input, and final design authority over the platform architecture.
- **Held sub-200ms inference latency at 99.5% uptime** across multi-region deployments, architecting Django + Docker services around PyTorch and TensorFlow pipelines.
- **Lifted annotation precision 15% and cut manual labeling effort 30%** by introducing GAN-based augmentation and transfer learning into the labeling pipeline.
- **Owned the client-facing technical relationship**, presenting quarterly product roadmaps directly to CTOs and engineering VPs at automotive, medical-imaging, and robotics customers.

EverestDX Inc.

Apr 2021 – Jan 2022
Hyderabad, India

Senior Engineer, Digital Platforms

- **Cut false-positive alerts by 25%** with an ML anomaly-detection platform for cloud infrastructure, and mapped 100+ cloud services into dependency dashboards used for incident triage.

Zen3 Info Solutions (now Tech Mahindra)

Aug 2019 – Mar 2021
Hyderabad, India

Software Engineer, AI Platforms

- **Shipped DATAMIME, a computer-vision platform producing ML-ready datasets, across 50+ data-collection campaigns** for automotive and ADAS clients in Europe and the Americas.
- **Trained and deployed OCR and object-detection models**, building the REST services that integrated them into the enterprise data platform.

INDEPENDENT PRODUCT ENGINEERING

Self-Directed Full-Stack Product & Infrastructure

2024 – Present
Hyderabad, India

> Sole architect and operator · **5 web services running in production 24/7** · **~1,600 automated tests** · test-driven, spec-first delivery

- **Design, ship and operate five production web services** on self-managed infrastructure spanning PostgreSQL, FastAPI, React PWAs, Cloudflare Tunnel edge routing, process supervision, nightly backups and uptime watchdogs. All live and publicly reachable.
- **Multi-tenant clinical management platform:** database-per-tenant isolation, offline-first React PWA, role-based access control, and a 736-test suite covering the full billing and records path.
- **Offline-first retail point-of-sale** engineered for intermittent connectivity, with an IndexedDB outbox, server-side reconciliation, row-level tenant isolation, and three-tier RBAC across 319 tests.
- **Organizational-memory RAG service:** pgvector hybrid retrieval with LLM synthesis over an internal knowledge base, exposed through FastAPI and a React client; 388 tests.
- **Systematic markets research stack:** event-driven ingestion, a backtesting harness with explicit look-ahead guards, and a paper-trading execution engine; 800+ tests.
- **Benchmarked local LLM inference across 8 models on a 64 GB M1 Max**, establishing that prefill scales with *active* parameters, not model size: a 3B-active mixture-of-experts prefilled 7.5× faster than a dense 27B. Of four serving engines, only one preserved prefix caching on hybrid-attention models, worth ~8s per turn.
- **Built evaluation harnesses that decide whether a model ships**. A 12-scenario discrimination gate caught a model scoring *worse* with a better prompt (8/12 → 6/12 given an explicit neutral rubric), i.e. following prompt emphasis over input data. A separate LLM decision layer was cut after 412 real decisions showed it degraded discrimination 18.5 points versus a rules baseline.

TECHNICAL SKILLS

AGENTIC AI	**AI agents & multi-agent orchestration**, **LangGraph (LangChain)**, tool use & function calling, Model Context Protocol (MCP), agentic planning, structured outputs, prompt & context engineering, Claude & OpenAI APIs
RETRIEVAL / RAG	**RAG architecture**, hybrid vector + keyword search, re-ranking, chunking strategy, embeddings, source attribution, Pinecone, ChromaDB, pgvector
LLMOPS & EVALS	**Evaluation harnesses**, **hallucination guardrails**, LLM-as-judge & rubric evals, retrieval-accuracy measurement, observability & tracing, latency & token-cost optimization, caching strategy, MLOps & CI/CD for AI systems
LOCAL INFERENCE	**MLX**, llama.cpp/GGUF runtimes, mixture-of-experts vs dense trade-offs, 4/8-bit quantization, prefix caching, KV-cache sizing, prefill vs decode profiling, benchmark selection
ML & VISION	**PyTorch**, scikit-learn, OpenCV, transfer learning, object detection, segmentation, OCR, NLP, GANs, TensorFlow
DATA & CLOUD	**Google Cloud Platform (GCP)**, **BigQuery**, **PostgreSQL** / pgvector, AWS, Docker, Cloudflare, DuckDB, SQLite, CI/CD · warehousing: Teradata
ENGINEERING	**Python**, TypeScript, SQL, Bash, FastAPI, Django, Flask, React, Streamlit, REST API design, test-driven development
LEADERSHIP	**Architecture & design authority**, technical direction, team leadership (15+), code-review culture, executive stakeholder management, AI strategy & roadmap, AI governance, mentoring

EDUCATION

B.Tech, Computer Science & Engineering
Presidency University, Bangalore
2015 – 2019

AWARDS & CERTIFICATIONS

- **Google Cloud Certified: Generative AI Leader** (2026)
- **Star of the Month & Quarter**, L&T (2023)
- **Python Expert (Level 5)**, HackerRank
- **Data Science Professional**, IBM