DEVENDRA REDDY YERVA

AI / GenAI Technical Lead | Multi-Agent Systems · RAG · Applied ML at Enterprise Scale

Hyderabad, India · +91 70133 13692 · dev225.reddys@gmail.com · linkedin.com/in/dev225x · devendra.me

7 yrs	3 days → 2 hrs	500+	15+	92%
PRODUCTION AI/ML	EXEC REVIEW CYCLE	ENTERPRISE USERS	ENGINEERS LED	RAG RETRIEVAL ACC.

PROFILE

Engineer and technical lead with **7 years building production AI systems** for Fortune 500 clients across telecom, automotive, and medical imaging. Currently the technical owner of a **multi-agent GenAI analytics platform** in Verizon's CFO organization — agent orchestration, NL-to-SQL over 50M+ row warehouses, and production RAG serving 500+ enterprise users. Previously led **15+ CV/ML engineers** taking an enterprise annotation platform from concept to eight international clients. Comfortable owning a system end-to-end: architecture, delivery, the executive conversation, and the on-call pager.

PROFESSIONAL EXPERIENCE

Verizon

Oct 2024 – Present
Hyderabad, India

Consultant — AI / GenAI | Business Transformation (CFO Organization)

> Technical owner of the CFO org's GenAI platform · **500+ enterprise users** · direct C-suite stakeholders · Claude Opus/Sonnet, GCP, BigQuery, Teradata

- **Cut executive business-case reviews from 3+ days to under 2 hours** by architecting BT OneView, a multi-agent GenAI analytics platform on GCP — sequential agent orchestration, tool-use routing, and structured-output parsing over the Claude API.
- **Opened 50M+ row Teradata and BigQuery datasets to non-technical stakeholders** through schema-aware NL-to-SQL pipelines, removing engineering from the query path entirely.
- **Returned 20+ analyst hours per week** with 8 autonomous overnight research agents that query the warehouse, synthesize findings, and deliver branded PDF/PPT briefings by morning — built with retry logic, error recovery, and supervisor-level quality gates.
- **Reached 92% retrieval accuracy with full source attribution** on a production RAG stack (Pinecone, ChromaDB) using hybrid vector + keyword search, chunk-size optimization, re-ranking, and hallucination guardrails.
- **Reduced incident MTTR by 40%** through multi-agent auto-triage — LLM agents performing classification, routing, and stakeholder notification without human dispatch.
- **Shipped 8 analytics dashboards to 500+ enterprise users**, backed by a three-tier DuckDB caching layer that holds sub-second response on warehouse-scale queries.
- **Set the AI automation roadmap alongside CFO-org leadership** in weekly strategy reviews — presenting ROI, defending measurement methodology, and shaping investment priorities across product, data engineering, and finance.

L&T Technology Services

Feb 2022 – Sep 2024
Hyderabad, India

Senior Engineer, Transportation Solutions — Technical Lead

> Led **15+ CV/ML engineers and data scientists** · **8 enterprise clients** across Europe & the Americas · 200+ daily users · 500K+ annotations/month

- **Took AnnotAI from concept to production** — a computer-vision annotation platform now serving 200+ daily users across 8 enterprise clients in automotive, medical imaging, and robotics at 500K+ annotations per month.
- **Led a team of 15+ engineers as technical lead** — owned architecture and design review, established sprint cadence, code-review standards, and CI/CD, and ran individual objective-setting and performance reviews.
- **Held sub-200ms inference latency at 99.5% uptime** across multi-region deployments, architecting Django + Docker services around PyTorch and TensorFlow pipelines.
- **Lifted annotation precision 15% and cut manual labeling effort 30%** by introducing GAN-based augmentation and transfer learning into the labeling pipeline.
- **Owned the client-facing technical relationship**, presenting quarterly product roadmaps directly to CTOs and engineering VPs at automotive, medical-imaging, and robotics customers.
- Awarded **Star of the Month & Star of the Quarter** (Aug–Nov 2023) for delivery leadership.

EverestDX Inc.

Apr 2021 – Jan 2022
Hyderabad, India

Senior Engineer — Digital Platforms

- **Cut false-positive alerts by 25%** with a machine-learning anomaly-detection platform for cloud infrastructure (Python, Pandas, scikit-learn).
- **Mapped 100+ cloud services** into interactive dependency dashboards built with US-based operations teams.
- **Automated weekly executive reporting**, returning roughly 10 hours per week to the team.

Zen3 Info Solutions (now Tech Mahindra)

Aug 2019 – Mar 2021
Hyderabad, India

Software Engineer — AI Platforms

- **Shipped DATAMIME across 50+ data-collection campaigns** for automotive and ADAS clients in Europe and the Americas — a computer-vision platform producing ML-ready datasets.
- **Trained and deployed OCR and object-detection models**, building the REST services that integrated them into the enterprise data platform.
- **Built custom NLP evaluation pipelines** to benchmark in-house speech models against commercial NLU baselines (Microsoft LUIS).

INDEPENDENT PRODUCT ENGINEERING

Self-Directed — Full-Stack Product & Infrastructure

2024 – Present
Outside working hours

> Sole architect and operator · **5 web services running in production 24/7** · **~1,600 automated tests** · test-driven, spec-first delivery

- **Design, ship, and operate five production web services** on self-managed infrastructure — PostgreSQL, FastAPI, React PWAs, edge routing via Cloudflare Tunnel, process supervision, nightly backup rotation, and uptime watchdogs. Every service is live and publicly reachable.
- **Multi-tenant clinical management platform** — database-per-tenant isolation, offline-first React PWA, role-based access control, and a 736-test suite covering the full billing and records path.
- **Offline-first retail point-of-sale** engineered for intermittent connectivity — IndexedDB outbox with server-side reconciliation, row-level tenant isolation, and three-tier RBAC across 319 tests.
- **Organizational-memory RAG service** — pgvector hybrid retrieval with LLM synthesis over an internal knowledge base, exposed through FastAPI and a React client; 388 tests.
- **Systematic markets research stack** — event-driven ingestion, a backtesting harness with explicit look-ahead guards, and a paper-trading execution engine; 800+ tests across 165K lines of Python.
- **Apply production discipline end-to-end** — test-driven development, conventional commits, spec-then-slice planning, and repo-tracked issue and decision ledgers on every project.

TECHNICAL SKILLS

GENAI & AGENTS	**Multi-agent orchestration**, **LangGraph**, tool use & function calling, Model Context Protocol (MCP), agentic planning, structured output parsing, prompt & context engineering, Claude & OpenAI APIs
RETRIEVAL / RAG	**RAG architecture**, hybrid vector + keyword search, re-ranking, chunking strategy, embeddings, source attribution, Pinecone, ChromaDB, pgvector
LLMOPS & EVAL	**Evaluation harnesses**, **hallucination guardrails**, retrieval-accuracy measurement, observability & tracing, latency & token-cost optimization, caching strategy, CI/CD for AI systems
ML & VISION	**PyTorch**, scikit-learn, OpenCV, transfer learning, object detection, segmentation, OCR, NLP, anomaly detection, generative models (GANs), TensorFlow
DATA & CLOUD	**GCP**, **BigQuery**, **PostgreSQL** / pgvector, AWS, Docker, Cloudflare, DuckDB, SQLite, Alembic, CI/CD · enterprise warehousing: Teradata
ENGINEERING	**Python**, TypeScript, SQL, Bash, FastAPI, Django, Flask, React, Streamlit, REST API design, test-driven development
LEADERSHIP	**Architecture & design authority**, technical direction, team leadership (15+), code-review culture, executive stakeholder management, roadmap definition, mentoring

EDUCATION

B.Tech, Computer Science & Engineering
Presidency University, Bangalore
2015 – 2019

AWARDS & CERTIFICATIONS

- **Star of the Month & Quarter** — L&T (2023)
- **Python — Expert (Level 5)**, HackerRank
- **Data Science Professional**, IBM